

Noname manuscript No.
(will be inserted by the editor)

Why Software Projects need Heroes (Lessons Learned from 1000+ Projects)

Suvodeep Majumder · Joymallya
Chakraborty · Amritanshu Agrawal ·
Tim Menzies

Received: date / Accepted: date

Abstract A “hero” project is one where 80% or more of the contributions are made by the 20% of the developers. Those developers are called “hero” developers. In the literature, heroes projects are deprecated since they might cause bottlenecks in development and communication. However, there is little empirical evidence on this matter. Further, recent studies show that such hero projects are very prevalent. Accordingly, this paper explores the effect of having heroes in project, from a code quality perspective by analyzing 1000+ open source GitHub projects. Based on the analysis, This study finds that (a) majority of the projects are hero projects; and (b) the commits from “hero developers” (who contribute most to the code) result in far fewer bugs than other developers. That is, contrary to the literature, heroes are standard and very useful part of modern open source projects.

Keywords Software Analytic, GitHub, Software Defects, Heroes, Social interaction, Code Interaction, Open Source Project

Suvodeep Majumder E-mail: smajumd3@ncsu.edu
Joymallya Chakraborty E-mail: jchakra@ncsu.edu
Amritanshu Agrawal E-mail: aagrawa8@ncsu.edu
Tim Menzies E-mail: timm@ieee.org

North Carolina State University
Raleigh, North Carolina
USA, 27695

1 Introduction

A “hero” project is one where 80% or more of the contributions come from 20% of the developers. Those developers are called “hero” developers. In the literature, hero projects are deprecated since, it is said, they are like bottlenecks that slow down the project development process and causes information loss [13, 18, 50, 77, 108, 39].

Recent studies have motivated a re-examination of the implications of heroes. In 2018, Agrawal et al. [2] studied 661 open source projects and 171 in-house proprietary projects. In that sample, over 89% of projects were hero-based¹. Only in the group of small open source projects (with under 15 core developers), non-hero projects were more prevalent.

To say the least, this widespread prevalence of heroes is at odds with established wisdom in the SE literature. The usual stance in the literature is to warn against heroes since they may become bottlenecks in development and communication [13, 18, 50, 77, 108, 91, 92, 24]. Hence, it is now an open and pressing issue to understand why so many projects are hero-based. To that end, this paper verifies the result which Agrawal et al. [2] found out. All of project data was recollected from scratch from double the number of open source projects (over 1000 projects) than used by Agrawal et al. In this study, heroes are those who participate in 80% (or more) of the total contribution in the project. At his ICSE’14 keynote, James Herbsleb stated that communication between developers is an important factor to find bugs when code interaction happens [47]. So, we decided not to look at the percent of code contribution but also at the percent of communication involvement. As a result, we will say that “hero” developer is a developer who participates in 80% or more of the code contribution or communications in a project. As shown below, the population of “heroes” defined in this way are an important segment of the development population (specifically, we will show that these “communication heroes” add far fewer bugs into software than anyone else).

Despite our different ways to recognize “heroes” and despite our much larger sample, this study comes to a similar conclusions as Agrawal et al.. This study finds majority of our projects contain heroes, which is very similar to the Agrawal et al.’s result. More importantly, this study can explain *why* heroes are more important. As shown below, our “hero” commit patterns (where “heroes” are those who interact most with other developers) are associated with dramatically fewer defects than the commits from non-heroes (who interact with fewer people). Hence we conjecture that heroes are so common since they write better code.

This is not the first paper to comment on the use of hero developers. For example, in 1975 Brooks [21] proposed basing programming teams around a small number of “chief programmers” (which we would call “heroes”) who are supported by a large number of support staff (Brooks’s analogy was the

¹ This text use “hero” for women and men since recent publications use it to denote admired people of all genders– see bit.ly/2UhJCek.

operating theater where one surgeon is supported by one or two anesthetists, several nurses, clerical staff, etc). The Agile Alliance [29] and Bach et al. [8] believed that heroes are the core ingredients in successful software projects saying "... the central issue is the human processor - the hero who steps up and solves the problems that lie between a need express and a need fulfilled." In 2002, Mockus et al. [75] analyzed Apache and Mozilla projects to show the presence of heroes in those projects and reported, surprisingly, their positive influence on the projects.

That said, this article is different from previous works because:

1. This study clearly demonstrate the benefits of hero-based development, which is contrary to much prior pessimism [13, 18, 50, 77, 108, 39].
2. Our conclusions come from over 1000+ projects, whereas prior work commented on heroes using data from just a handful of projects [9, 65, 112, 35, 19, 16, 11, 10, 62, 54, 67].
3. Our conclusions come from very recent projects instead of decades-old data [52, 43, 54, 57, 11, 111].
4. This study shows curves that precisely illustrate the effects on code quality for different levels of communication. This is different to prior works that only offered general qualitative principles [106, 33, 42, 48, 27, 26].
5. As discussed in Section 2.2, this paper makes its conclusions using more metrics than prior work. Not only do we observe an effect (using *process* and *resource* metrics) to report the frequency of developer contribution, but we also report the consequence of that effect (by joining to *product metrics* to reveal software quality).
6. Instead of just reporting an effect (that heroes are common, as done by Agrawal et al. [2]) this study can explain that effect (heroes are those that communicate more and that communication leads to fewer bugs).
7. As a service to other researchers, all the scripts and data of this study can be downloaded from tiny.cc/git_mine.

The practical implication of this research is as follows:

- Our results demand a change in the way we develop new technologies or modern software projects. Specifically, given the prominence and importance of heroes, future work could usefully explore methods to streamline the communication between a *very large population* of developers and a *very small number* of heroes that are critical for high quality software.

Before beginning, we make some definitional points. Firstly, when we say 1000+ projects, that is shorthand for the following. Our results used the intersection of two graphs of *code interaction graph* (of who writes what and whose code) with *social interaction graph* (who discusses what issue and with whom) from 1037 projects. Secondly, by code interaction graphs and social interaction graphs, we mean the following. Each graph has own nodes and edges $\{N, E\}$. For code interaction graphs:

- Individual developers have their own node N_c ;




Fig. 1 An example of social interaction graph generated from our data. The number of nodes equals the number of unique people participating in issue conversation. The existence and width of each edge represents the frequency of conversation between pairs of developers. Hero programmers are those nodes which have very high node degree (i.e. who have participated in lot of unique conversations). Note that, in this example data, these hero programmers are few in number.

- The edge E_c connects two nodes and indicates if ever one developer has changed another developer's code. W_c denotes how much one developer has changed another developer's code.

For social interaction graphs like Figure 1:

- A node N_s is created for each individual who has created or commented on an issue.
- An edge E_s indicates a communication between two individuals (as recorded in the issue tracking system). If this happens N times then the weight $W_s = N$.

Thirdly, we have defined heroes based on code contribution and communication. From the “Code interaction graph”, the developers who contribute more than 80% are hero contributors and in the “Social interaction graph”, the developers who are making 80% of the communication are hero communicators. Both are “hero developers” for us. As we show in §4, both these definitions of heroes are insightful since more can be predicted about a project using *both* definitions that it either is applied separately.

The rest of the paper is organized into the following sections. Section 2 provides background information that directly relates to our research questions, in addition to laying out the motivation behind our work. Section 3.1 explains the data collection process and Section 3.2, a detailed description of our experimental setup and data is given, along with our performance criteria for evaluation is presented. It is followed by Section 4 the results of the experiments and answers to some research questions. Section 5 discusses their implication of our finding. It is followed by Section 6, which discusses threats to validity. Finally Section 7 concludes the paper.

2 Background And Prior Work

2.1 Heroism in Software Development

Heroism in software development is a widely studied topic. Various researchers have found the presence of heroes in software projects. For example:

- Peterson analyzed the software development process on GitHub and found out a pattern that most development is done by a small group of developers [83]. He stated that for most of the GitHub projects, 95-100% commits come from very few developers.
- In 2002, Koch et al. [59] studied the GNOME project and showed the presence of heroes through out the project history. They conjectured (without proof) that the small number of hero developers may allow easy communication and collaboration. Interestingly, they also showed there is no relation between developer’s time in the project and being a hero developer.
- In 2005, Krishnamurthy [61] studied 100 open source projects to find that a few individuals are responsible for the main contribution of the project in most of the cases.
- In 2006 and 2009, Robles et al. [93, 92] explored in their research the presence and evolution of heroes in open source software community.
- In 2018, Agarwal et al. [2] stated that hero projects are very common. In fact, as software projects grow in size, nearly all projects become hero projects.

Most prior researchers deprecate heroism in software projects. They argue that

- Having most of the work being dependent on a small number of heroes can become a bottleneck that slows down project development [13, 77, 50, 18, 108].

- In the case of hero projects, there is less collaboration between team members since there are few active team members. So, heroes are affecting the collaboration which is essential [7, 104].

This second point is problematic since, in the literature, studies that analyze distributed software development on social coding platforms like GitHub and Bitbucket [34, 31] remark on how social collaborations can reduce the cost and efforts of software development without degrading the quality of software. Distributed coding effort is beneficial for agile community-based programming practices which can in turn have higher customer satisfaction, lower defect rates, and faster development times [76, 87]. Customer satisfaction, it is argued, is increased when faster development leads to:

- Increasing the number of issues/bugs/enhancements being resolved [75, 51, 17, 6, 44, 90].
- Lowering the issues/bugs/enhancements resolution times [51].

Even more specifically, as to issues related to heroes, Bier et al. warn when project becomes complicated, it is always better to have a community of experts rather than having very few hero developers [13]. Williams et al. have shown that hero programmers are often responsible for poorly documented software system as they remain more busy in coding rather than writing code related documents [50]. Also, Wood et al. [108] caution that heroes are often code-focused but software development needs workers acting as more than just coders (testers, documentation authors, user-experience analysts).

Our summary of the above is as follows: with only isolated exceptions, most of the literature deprecates heroes. Yet as discussed in the introduction, many studies indicate that heroic projects are quite common. This mismatch between established theory and a widely observed empirical effect prompted the analysis discussed in this paper.

2.2 Software Quality Metrics

Table 1 offers some context for this research. While this paper product, process, and personnel metrics for 1000+ projects, most of the papers explore far fewer projects using a narrower range of metrics. For example, as shown by the *Number of Projects* column in Table 1, our sample size (1000+ projects) is orders of magnitude larger than the typical paper in this arena.

This table was generated as follows. Firstly, using Google Scholar we searched for “(software heroes) or ((software metrics) and (code quality))”. Secondly, for papers more than two years old, we pruned “non-influential papers” which we define has having less than ten citations per year. Thirdly, we read the papers to determine what kind of metrics they used. When presenting these results (in Table 1), hero-related publications have a blue background (bold text also) while rows colored in gray denote hero-related publication that offer no metrics in support of their arguments.

Table 1 Some results from Google Scholar query (*software heroes*) or ((*software metrics*) and (*code quality*)). Hero-related publications have a color background. Rows colored in gray denote hero-related publication that offer no metrics in support of their arguments.

Ref	Year	cites	#projects analyzed	Uses Product Metrics?	Uses Process Metric?	Uses Personnel Metrics?
[9]	1996	1994	8	✓		
[75]	2002	1961	2		✓	✓
[65]	1993	1268	2	✓		
[20]	2000	779	1	✓		
[79]	2006	772	5	✓	✓	
[97]	2002	711	1	✓	✓	
[113]	2007	636	3	✓	✓	
[18]	2006	667	0		✓	
[81]	2005	622	2	✓	✓	
[112]	2009	466	12	✓	✓	
[61]	2002	466	100		✓	
[35]	2001	445	1	✓		
[37]	2001	406	1	✓		
[25]	2000	400	1	✓		
[36]	2008	398	4	✓		
[19]	1999	346	1	✓		
[59]	2002	305	1		✓	
[4]	1999	300	3	✓		
[104]	2007	298	0		✓	
[95]	2009	271	1	✓	✓	
[73]	2010	256	10	✓		
[96]	2011	256	17	✓	✓	
[16]	2011	233	2		✓	✓
[55]	2010	229	38	✓		
[38]	2004	223	30		✓	
[84]	2008	223	1	✓	✓	✓
[72]	2008	218	1	✓	✓	✓
[107]	2009	197	1		✓	✓
[3]	2005	186	SourceForge			✓
[15]	2009	177		6	✓	✓
[14]	1998	172		2	✓	
[103]	2008	163		3	✓	✓
[11]	2012	163		11	✓	✓
[80]	2014	159		9	✓	✓
[58]	2006	131		1	✓	✓
[45]	2015	106		10	✓	✓
[70]	2012	103		905,470		✓
[88]	2008	102		5	✓	
[92]	2006	99	21	✓		
[71]	2016	92		3		✓
[111]	2014	87		1,398	✓	
[68]	2002	85		39,000		✓
[63]	2015	85		0	✓	✓
[69]	2015	76		18	✓	✓
[78]	2013	68	0		✓	
[93]	2009	65	1		✓	✓
[66]	2014	61	GitHub			✓
[82]	2010	59		6	✓	✓
[17]	2013	58	100,000	✓	✓	
[1]	2009	54	1		✓	✓
[53]	2011	48		2	✓	✓
[10]	2013	37		3	✓	✓
[108]	2005	36		0		
[30]	2010	30		2	✓	
[12]	2011	27		2		✓
[51]	2014	24	2,000		✓	
[100]	2007	22		4	✓	
[110]	2016	19		235,000	✓	✓
[83]	2013	14	1,000		✓	
[86]	2015	12		1	✓	✓
[62]	2017	11		10	✓	✓
[54]	2018	11		15	✓	✓
[2]	2018	6	832		✓	
[43]	2017	5		12	✓	
[13]	2011	3		0		
[52]	2018	2		4	✓	✓
[50]	2002	2		0		
[77]	2012	2		0		
[89]	2018	0		5	✓	
[85]	2018	0		1	✓	
[32]	2018	0		2		✓
[67]	2018	0		1	✓	
[102]	2017	0		50		✓
[98]	2018	0		0	✓	✓

Table 1 also shows that most papers do not use a wide range of software metrics. Xenos [109] distinguishes software metrics as follows. *Product metrics* are metrics that are directly related to the product itself, such as code statements, delivered executable, manuals, and strive to measure product quality, or attributes of the product that can be related to product quality. *Process metrics* focus on the process of software development and measure process characteristics, aiming to detect problems or to push forward successful practices. Lastly, *personnel metrics* (a.k.a. *resource metrics*) are those related to the resources required for software development and their performance. The capability, experience of each programmer and communication among all the programmers are related to product quality [106, 33, 27, 26]. In our work:

- Code interaction graph is a process metric;
- Social interaction graph is a personnel metric;
- Defect counts are product metrics.

(Aside: In this text we have used “resource” and “peronnel” interchangeably since, according to Center for Systems and Software Engineering, [109] resource metrics relating to programmer quality or communication related metrics are also called *personnel metrics*.)

This paper explores all three kinds of metrics and applies the combination to exploring the effects of heroism on software development. There are many previous studies that explore one or two of these types of metrics. Fig 2 summarizes Table 1 and shows that, in that sample, very few papers in software metrics and code quality combine insights from product and process and personnel metrics. To the best of our knowledge, this is the first paper in this arena to discuss heroism using product and process and personnel metrics.

Having worked with that data, we think we know why other publications do not report results using a wide range of metrics. Such reports require extensive and elaborate queries. The analysis of this paper required months of struggling with the GitHub API (and its queries/hour limits), followed by much scripting, followed by many tedious manual checks that our automatic tools were behaving sensibly. In all, we estimate that this paper required nine weeks of coding (40 hours per week) to join across process and product and personnel metrics.




Fig. 2 Summary of Table 1

2.3 Herbsleb Hypothesis (and Analogs)

One way to view this paper is as a check of the *Hersleb hypothesis* [47]. At his ICSE'14 keynote, James Hersleb defined coding to be a socio-technical process where code and humans interact. According to the Hersleb hypothesis [47], the following anti-pattern is a strong predictor for defects:

- If two code sections communicate...
- But the programmers of those two sections do not...
- Then that code section is more likely to be buggy.

To say that another way:

Coding is a social process and better code arises from better social interactions.

Many other researchers offer conclusions analogous to the Herbsleb hypothesis. Developer communication/interaction is often cited as one of the most important factor for a successful software development [28,60,46]. Many researchers have shown that successful communication between developers and adequate knowledge about the system plays a key role in successful software development [99,41,64]. As reported as early as 1975 in Brooks et al. text “The Mythical Man Month” [23], communication failure can lead to coordination problem, lack of system knowledge in the projects as discussed by Brooks et al. in the Mythical Man-Month.

The usual response to the above argument is to improve communication by “smoothing it out”, i.e. by deprecating heroes since, it is argued, that encourages more communication across an entire project [13,18,50,77,108]. The premise of “smoothing it out” is that heroes are bad and should be deprecated. This paper tries to verify whether or not this premise holds true for open source GitHub projects or not.

3 Methodology

3.1 Data Collection

Figure 3 summarizes the Github data we used for this study. To understand this figure, we offer the following definitions:

- *Release*: (based on Git tags) mark a specific point in the repository’s history. Number of releases defines different versions published, which signifies considerable amount of changes done between each version.
- *Duration*: length of project from its inception to current date or project archive date. It signifies how long a project has been running and in active development phase.
- *Stars*: signifies number of people liking a project or use them as bookmarks so they can follow what’s going on with the project later.




Fig. 3 Distribution of projects depending on Number of Releases, Duration of Project, Number of Stars, Forks, Watchers and Developers. Box plots show the min to max range. Central boxes show the 25th, 50th, 75th percentiles.

- *Forks*: A fork is a copy of a repository. Forking a repository allows user to freely experiment with changes without affecting the original project. This signifies how people are interested in the repository and actively thinking of modification of the original version.
- *Watcher*: Watchers are GitHub users who have asked to be notified of activity in a repository, but have not become collaborators. This is a representative of people actively monitoring projects, because of possible interest or dependency.
- *Developer*: Developers are the contributors to a project, who work on some code, and submit the code using commit to the codebase. The number of developers signifies the interest of developers in actively participating in the project and volume of the work.

Figure 4 shows that the projects we chose for our experiment comprise different languages. Note that we did not use all the data in Github. GitHub has over 100 million repositories as of May, 2019 so we only use data from the “GitHub showcase project” list. Many of these projects contain very short development cycles; are used for personal use; and are not related to software development. Such projects may bias research findings. To mitigate that, we filter out projects using the standard “sanity checks” recommended in the literature [56, 78]:

- *Collaboration*: refers to the number of pull requests. This is indicative of how many other peripheral developers work on this project. We required all projects to have at least one pull request.
- *Commits*: The project must contain more than 20 commits.
- *Duration*: The project must contain software development activity of at least 50 weeks.
- *Issues*: The project must contain more than 10 issues.
- *Personal Purpose*: The project must not be used and maintained by one person. The project must have at least eight contributors.
- *Software Development*: The project must only be a placeholder for software development source code.




Fig. 4 Distribution of projects depending on languages. Many projects use combinations of languages to achieve their results. Here, we show majority language used in the project.

- *Project Documentation Followed:* The projects should follow proper documentation standard to log proper commit comment and issue events to allow commit issue linkage.
- *Social network validation:* The Social Network that is being built should have at least 8 connected nodes in both the communication and code interaction graph (this point is discussed further in 3.2.2 and 3.2.3).

For each of the selected projects, the study recreates all the committed files to identify code changes in each commit file and identifies developers using the GitHub API, then downloads issue comments and events for a particular project, and uses the `git log` command to mine the git commits added to the project throughout the project lifetime. Using the information from each commit message, this study uses keyword based identifier [94, 49, 101] to label commits as buggy commit or not by identifying commits which were used to fix some bugs in the code and then identifies the last commit which introduced the bug. This commit is labeled as *buggy commit*.

3.2 Metric Extraction

3.2.1 Process Metrics

Recall that the developer code interaction graph records who touched what and whose code, where a developer is defined as a person who has ever committed any code into the codebase. We create that graph as follows:

- Project commits were extracted from each branch in the git history.
- Commits are extracted from the git log and stored in a file system.
- To access the file changes in each commit we recreate the files that were modified in each commit by (a) continuously moving the `git head` chronologically on each branch. Changes were then identified using `git diff` on two consecutive git commits.
- The graph is created by going through each commit and adding a node for the committer. Then we use `git blame` on the lines changed to find previous commits following a similar process of SZZ algorithm [105]. We identify all the developers of the commits from `git blame` and add them as a node as well.
- After the nodes are created, directed edges were drawn between the developer who changed the code, to whose code was changed. Those edges were weighted by the change size between the developers.

3.2.2 Personnel Metrics

Recall that the developer social interaction graph records who talked to each other via issue comments. We create that graph as follows.

- A node is created for the person who has created the issue, then another set of nodes are created for each person who has commented on the issue. So essentially in Social interaction graph each node in the graph is any person (developer or non-developer) ever created an issue or commented on an issue.
- The nodes are connected by directed edges, which are created by (a) connecting from the person who has created the issue to all the persons who have commented in that issue and (b) creating edges from the person commenting on the issue to all other persons who have commented on the issue, including the person who has created the issue.
- The edges are weighted by the number of comments by that person.
- The weights are updated using the entire history of the projects. The creation and weight update is similar to Figure 5.

3.2.3 Product Metrics

This study explores the effects of social and code communication to assess code quality, by measuring the percentage of buggy commits introduced by




Fig. 5 Example of creating a social interaction graph between four GitHub developers. **Step 1 (LHS):** Ananya, Maria, Chen and David are four developers in a GitHub project. **Step 2:** Ananya creates one issue where Maria, Chen and David comment. So, we join Ananya-Maria, Ananya-Chen, Ananya-David with edge of weight 1. **Step 3:** A new developer Vadim comes. **Step 4 (RHS):** Vadim creates one new issue where Ananya and David comments. So, two new edges are introduced - (Ananya-Vadim(1), David-Vadim(1)). Now we iterate for each developer, so all of them become connected and lastly, weight of Ananya-David increases to 2.

developers (hero and non-hero developers), but in order to do so we do need to identify the commits that introduced the bug in the code from the historic project data. This is a challenging task since there is no direct way to find the commits or the person who is responsible for the bug/issue introduction. Hence, our scripts proceed as follows:

- Starting with all the commits from `git log`, we identify the commit messages.
- Next, to use the commit messages for labeling, we apply natural language processing [49, 94] (to do stemming, lemmatization and lowercase conversion to normalize the commit messages).
- Then to identify commit messages which are representation of bug/issue fixing commits, a list of words and phrases is extracted from previous studies of 1000+ projects (Open Source and Enterprise). The system checked for these words and phrases in the commit messages and if found, it marks these as commits which fixed some bugs.
- To perform sanity check, 5% of the commits was manually verified by 7 graduate students using random sampling from different projects. Disagreements between manual labeling and keyword based labeling was further verified and keywords were added or removed to improve performance.

- These labeled commits were then processed to extract the file changes as the process mentioned in process metrics section 3.2.1.
- Finally, `git blame` is used to go back in the git history to identify a responsible commit where each line was created or changed last time.

By this process, commits that were responsible for introduction of the bugs in the system/project can be found. We label these commits as “buggy commits” and label the author of the commit as the “person responsible” for introducing the bug.

3.2.4 Final Feature Extraction

To assess the prevalence of heroes in the software projects, we joined across all the metrics shown above. Specifically, using the two graphs, we calculated the node degree (number of edges touching a vertex) of the graphs (and note that vertices with higher degree represent more communication or interaction).

For the sake of completeness, we varied our threshold of “hero” to someone belonging to 80%, 85%, 90% and 95% of the communication. In our studies, top contributors (or heroes) and non-heroes were defined as :

$$\text{Node Degree of } N_i = D(N_i) = \sum_{j=1}^n a_{ij} \quad (1)$$

$$\text{Hero} = \text{Rank}(D(N_i)) > \frac{P}{100} * (N + 1) \quad (2)$$

$$\text{Non-Hero} = \text{Rank}(D(N_i)) < \frac{P}{100} * (N + 1) \quad (3)$$

where:

N	Number of Developers
P	80,85,90 and 95 Percentile
Rank()	The percentile rank of a score is the percentage of scores in its frequency distribution that is equal to or lower than it.
a	Adjacency matrix for the graph where $a_{ij} > 0$ denotes a connection

Using these data and by applying the hero definition from formula (2) and (3) (look at the top 20%, 15%, 10% and 5%), we can find the developers who are responsible for 80%, 85%, 90% and 95% of the work. We use this to categorize the developers into 2 groups:

- The *hero developers*; i.e. the core group of the developers of a certain project who make regular changes in the codebase. In this study this is represented by the developers whose node degree is above 80, 85, 90 and 95th percentile of the node degree (developers’ communication and code interaction of the system graph).

- The *non-hero developers* are all other developers; i.e. developers associated with nodes with a degree below the respective threshold percentile.

Following this for each selected projects, we merge the product data collected in section 3.2.3 and section 3.2.1 to find each developer's code contribution according to the code interaction graph. Similarly process is followed for section 3.2.2 and section 3.2.1 in the social interaction graph. Using the above mentioned data, we can validate the code and social contribution of each developer along with their bug introduction percentage. This information will help us to answer the research questions asked in this study.

4 Results

Our results are structured around three research questions:

RQ1: How common are hero projects?

RQ2: What impact does heroism have on code quality?

RQ3: Does the results support Herbsleb Hypothesis?

RQ1: *How common are hero projects?*

Figure 6 shows the result for number of projects marked as hero and non-hero when we vary the threshold from “hero-ness” from 80 to 95%. The clear conclusion from this figure is that the phenomenon that we have defined as “hero” is very common. In fact, the phenomenon may be more pronounced than previously reported. Even when we require heroes to be involved in 95% of the communication (which is a large amount), we find that that majority of the projects studied here exhibit “hero-ness”. That is:

Result: Hero projects are overwhelmingly present in open source software community. That is, in the usual case, there are very few people in each project responsible for majority of the work.

RQ2: *What impact does heroism have on code quality?*

RQ2 explores effect heroism have on code quality. In this study, we created the developer social interaction graph and developer code interaction graph, then identified the developer responsible for introducing those bugs into the codebase. Then we find the percentage of buggy commits introduced by those developers by checking (a) the number of buggy commits introduced by those developers and (b) their number of total commits.

Fig 7 and Fig 8 shows the comparison between the performance of hero and non-hero developers. In those figures:

- The x-axis is different projects used in this study.




Fig. 6 RQ1 Result: Distribution of hero projects vs non-hero projects based on hero threshold being 80%, 85%, 90% and 95% respectively. Here threshold being 80% means in that project 80% code is done by less than 20% of developers

- The y-axis represents the median of the bug introduction percentage for all hero and non-hero developers for each project respectively.

Here projects are sorted by the number of non-hero developers. In those charts we note that:

- There exists a large number of non-heroes who always produce buggy commits, 100% of the time (evidence: the flat right-hand-side regions of the non-hero plots in both figures). That population size of “always buggy” is around a third in Fig 7 and a fourth in Fig 8.
- To say the least, heroes nearly always have fewer buggy commits than non-heroes. The 25th, 50th, 75th percentiles for both groups are shown in Table 2. This table clearly shows why heroes are so prevalent, they generate commits that are dramatically less buggy than non-heroes, regardless of the size of the project.

Table 2 The table summarizes of Fig 7, Fig 8 and stratifies the data according to 25th, 50th and 75th percentile of the developers.

Metric	Group	Percentile		
		25th	50th	75th
Code Interaction	Hero	0.52	0.58	0.53
	Non-Hero	0.67	0.75	1.0
	Ratio	1.3	1.3	1.9
Social Interaction	Hero	0.44	0.5	0.5
	Non-Hero	0.67	0.75	0.67
	Ratio	1.5	1.5	1.3




Fig. 7 Code interaction graph results for RQ2: Percentage of bugs introduced by hero and non-hero developers from developer code interaction perspective in Hero Projects .

The other thing to note from Fig 7 and Fig 8 is that they are nearly identical. That is, no matter how we define “hero”, we reach the same conclusions. Hence we say -

Result: In modern software projects, reflecting on who writes most of the code is just as insightful as reflecting on who participates in most of the discussion about the code.




Fig. 8 Social interaction graph results for RQ2: Percentage of bugs introduced by hero and non-hero developers from developer social interaction perspective in Hero Projects.

RQ3: *Do the results support Herbsleb Hypothesis?*

In this research question, we explored the Herbsleb hypothesis [47] from section 2; i.e. does lack of communication between developers predict for bugs in the code? In order to do that for 1,037 projects, we discretized the developers into 3 groups (i.e. High, Medium and Low) based on their code contribution (Code Node Degree) and social communication frequency (Social Node Degree). In Table 3, group HH (High,High) represents the developers who have

high high code contribution and social communication frequency and the value in the cell is median bug introduction percentage, while group LL (Low,Low) represents the developers who have low code contribution and low social communication frequency.

rank	group (code,social)	median	IQR	
1	L,H	30	29	●————
1	M,H	38	28	●————
2	L,M	38	31	●————
2	H,H	42	18	●————
2	M,M	46	21	●————
2	H,M	46	21	●————
2	H,L	48	30	●————
3	M,L	52	29	●————
4	L,L	67	55	●————

Table 3 This figure shows the result of statistical significance test and an effect size test on 9 different groups used to study the Herbsleb hypothesis. In this figure the “group” column represents the 9 different groups in this research question, where the first character represents the code node degree, while the later is social node degree.

Table 3 shows the result of statistical test performed on the 9 different groups representing different frequency and volume of communication. The “rank” column of that table shows a statistical analysis where one row has a higher rank than the previous one only if the two rows are

- Statistically *significant different*. For this test, we used the the Scott-Knot method recommended by Angelis and Mittas at TSE’13 [74].
- And that difference is *not a small effect*. For this effect size test, we used the the A12 test recommended by Angelis and Briand at ICSE’11 [5]

These statistical tests were selected since they were non-parametric; i.e. they do not assume normal distributions.

To summarize the effect reported in Table 3, we need to look at the difference between the highest and lowest ranks:

- In the groups with highest defects and highest rank, the social and coding groups are both low. That, is non-hero-ness (for both code and social interaction) is associated with worst quality code.
- In the groups with lowest defects and lowest rank, the social group is always high while the code groups are either low or medium. That is, extensive social interaction even with low to medium code interaction is associated with best quality code.

From the above, we can say that:

- Prior definitions of “hero” based just on code interaction need now to be augmented. As shown in Table 3, social hero-ness is much more predictive for bugs than merely reflecting on code hero-ness
- Also, these results support the Herbsleb hypothesis, which suggests communication between developers is an important factor and lesser social communication can lead to more bugs.

This finding leads to the following conjecture (to be explored in future work): the best way to reduce communication overhead and to decrease defects is to *centralize the communicators*. In our data, commits with lower defects come from the small number of hero developers who have learned how to talk to more people. Hence, we would encourage more research into better methods for rapid, high-volume, communication in a one-to-many setting (where the “one” is the hero and the “many” are everyone else). In summary, we can say

Result: The Herbsleb hypothesis holds true for open source software projects. More research should be performed to find better methods for rapid, high-volume, communication in a one-to-many setting.

5 Discussion

One strange feature of our results is that what is old is now new. Our results (that heroes are important) echo a decades old concept. In 1975, Fred Brooks wrote of “surgical teams” and the “chief programmer” [22]. He argued that:

- Much as a surgical team during surgery is led by one surgeon performing the most critical work, while directing the team to assist with less critical parts.
- Similarly, software projects should be led by one “chief programmer” to develop critical system components while the rest of a team provides what is needed at the right time.

Brooks conjecture that “good” programmers are generally much more productive than mediocre ones. This can be seen in the results that hero programmers are much more productive and less likely to introduce bugs into the codebase. Heroes are born when developers become so skilled at what they do, that they assume a central position in a project. In our view, organizations need to acknowledge their dependency on such heroes, perhaps altering their human resource policies and manage these people more efficiently by retaining them.

6 Threats to Validity

6.1 Sampling Bias

Our conclusions are based on 1000+ open source GitHub projects that started this analysis. It is possible that different initial projects would have lead to different conclusions. That said, our initial sample is very large so we have some confidence that this sample represents an interesting range of projects.

6.2 Evaluation Bias

In RQ1,RQ2 and RQ3, we said that heroes are prevalent and responsible for far less bug introduction than non-hero developers. It is possible that, using other metrics² then there may well be a difference in these different kinds of projects. But measuring people resources only by how fast releases are done or issues are fixed may not be a good indicator of measuring affects of having heroes in team. This is a matter that needs to be explored in future research.

Another evaluation bias as we report cumulative statistics of lift curves where other papers reported precision and recall. The research in this field is not mature enough yet for us to say that the best way to represent results is one way versus another. Here we decided to use lift curves since, if we'd used precision and recall, we had to repeat that analysis at multiple points of the lift curve. We find our current lift curves are a succinct way to represent our results.

6.3 Construct Validity

At various places in this report, we made engineering decisions about (e.g.) team size; and (e.g.) what constitutes a “hero” project. While those decisions were made using advice from the literature (e.g. [40]), we acknowledge that other constructs might lead to different conclusions.

That said, while we cannot prove that all of our constructs are in any sense “optimum”, the results of Table 3 suggest that our new definition of social hero-ness can be more informative than constructs used previously in the literature (that defined “hero” only in terms of code interaction).

Another issue about our construct validity is that we have relied on a natural language processor to analyze commit messages to mark them as buggy commits. These commit messages are created by the developers and may or may not contain proper indication of if they were used to fix some bugs. There is also a possibility that the team of that project might be using different syntax to enter in commit messages.

Yet another threat to construct validity is that we did not consider the different roles of the developers. We had trouble extracting that information from our data source, we found that people have multiple roles particularly our heroes who would often step in and assist in multiple activities. Nevertheless the exploration of different roles would be an interesting study.

6.4 External Validity

Previously Agrawal et al. were able to comment on the effects of heroes in open and closed source projects. That research group was fortunate enough to work on-site at a large open source software company. We were not as fortunate as

² E.g. do heroes reduce productivity by becoming bottleneck

them. We therefore acknowledge our findings (from open source projects) may not be the same for closed source projects.

Similarly we have used GitHub issues and comments to create the communication graph, It is possible that the communication was not made using these online forums and was done with some other medium. To reduce the impact of this problem, we did take precautionary step to (e.g.,) include various tag identifiers of bug fixing commits, did some spot check on projects regarding communication etc.

Our conclusion shows that almost all (when experimenting with 80%, 85%, 90% threshold) of our sample projects are hero dominated. In case of large size public GitHub projects, there are official administrators and maintainers who are responsible for issue labelling or assigning. So, they frequently comment on all of the issues but though they are not active developers. These people should not be considered as hero developers. Finding these people needs manual inspection which is not possible for 1000+ projects. We decide to put it as a limitation of our study as we deal with a huge number of projects.

We do not isolate hero projects and non-hero projects and look into them separately because there are very few non-hero projects and also there are a lot of developers who work in a large number of projects (some of them are hero projects and some of them are not).

7 Conclusion

The established wisdom in the literature is to deprecate “heroes”, i.e., a small percentage of the staff who are responsible for most of the progress on a project. But, based on a study of 1000+ open source GitHub projects, we assert:

- Overwhelmingly, most projects are hero projects.
- Hero developers are far less likely to introduce bugs into the codebase than their non-hero counterparts. Thus having heroes in projects significantly affects the code quality.

Our empirical results call for a revision of a long-held truism in software engineering. Software heroes are far more common and valuable than suggested by the literature, particularly from code quality perspective. Organizations should reflect on better ways to find and retain more of these software heroes.

More generally, we would comment that it is time to reflect more on long-held truisms in our field. Heroes are widely deprecated in the literature, yet empirically they are quite beneficial. What other statements in the literature need to be reviewed and revised?

8 Acknowledgements

This work was partially funded by NSF Grant #1908762.

References

1. Roberto Abreu and Rahul Premraj. How developer communication frequency relates to bug introducing changes. In *Proceedings of the Joint International and Annual ERCIM Workshops on Principles of Software Evolution (IWPSE) and Software Evolution (Evol) Workshops*, IWPSE-Evol '09, pages 153–158, New York, NY, USA, 2009. ACM.
2. Amritanshu Agrawal, Akond Rahman, Rahul Krishna, Alexander Sobran, and Tim Menzies. We don't need another hero? *Proceedings of the 40th International Conference on Software Engineering Software Engineering in Practice - ICSE-SEIP '18*, 2018.
3. and, S. Christley, and G. Madey. A topological analysis of the open souce software development community. In *Proceedings of the 38th Annual Hawaii International Conference on System Sciences*, pages 198a–198a, Jan 2005.
4. and and. An empirical study on object-oriented metrics. In *Proceedings Sixth International Software Metrics Symposium (Cat. No.PR00403)*, pages 242–249, Nov 1999.
5. A. Arcuri and L. Briand. A practical guide for using statistical tests to assess randomized algorithms in software engineering. In *ICSE'11*, pages 1–10, 2011.
6. Dimitrios Athanasiou, Ariadi Nugroho, Joost Visser, and Andy Zaidman. Test code quality and its relation to issue handling performance. *IEEE Transactions on Software Engineering*, 40(11):1100–1125, 2014.
7. L. Augustin, D. Bressler, and G. Smith. Accelerating software development through collaboration. In *Proceedings of the 24th International Conference on Software Engineering. ICSE 2002*, pages 559–563, May 2002.
8. James Bach. Enough about process: what we need are heroes. *IEEE Software*, 12(2):96–98, 1995.
9. V. R. Basili, L. C. Briand, and W. L. Melo. A validation of object-oriented design metrics as quality indicators. *IEEE Transactions on Software Engineering*, 22(10):751–761, Oct 1996.
10. Robert M. Bell, Thomas J. Ostrand, and Elaine J. Weyuker. The limited impact of individual developer data on software defect prediction. *Empirical Software Engineering*, 18(3):478–505, Jun 2013.
11. Pamela Bhattacharya, Marios Iliofotou, Iulian Neamtiu, and Michalis Faloutsos. Graph-based analysis and prediction for software evolution. In *Proceedings of the 34th International Conference on Software Engineering*, ICSE '12, pages 419–429, Piscataway, NJ, USA, 2012. IEEE Press.
12. Serdar Biçer, Ayşe Başar Bener, and Bora Çağlayan. Defect prediction using social network analysis on issue repositories. In *Proceedings of the 2011 International Conference on Software and Systems Process*, ICSSP '11, pages 63–71, New York, NY, USA, 2011. ACM.
13. Norman Bier, Marsha Lovett, and Robert Seacord. An online learning approach to information systems security education. In *Proceedings of the 15th Colloquium for Information Systems Security Education*, 2011.
14. A. B. Binkley and S. R. Schach. Validation of the coupling dependency metric as a predictor of run-time failures and maintenance measures. In *Proceedings of the 20th International Conference on Software Engineering*, pages 452–455, April 1998.
15. C. Bird, N. Nagappan, H. Gall, B. Murphy, and P. Devanbu. Putting it all together: Using socio-technical networks to predict failures. In *2009 20th International Symposium on Software Reliability Engineering*, pages 109–119, Nov 2009.
16. Christian Bird, Nachiappan Nagappan, Brendan Murphy, Harald Gall, and Premkumar Devanbu. Don't touch my code!: Examining the effects of ownership on software quality. In *Proceedings of the 19th ACM SIGSOFT Symposium and the 13th European Conference on Foundations of Software Engineering*, ESEC/FSE '11, pages 4–14, New York, NY, USA, 2011. ACM.
17. Tegawendé F Bissyandé, David Lo, Lingxiao Jiang, Laurent Réveillere, Jacques Klein, and Yves Le Traon. Got issues? who cares about it? a large scale investigation of issue trackers from github. In *Software Reliability Engineering (ISSRE), 2013 IEEE 24th International Symposium on*, pages 188–197. IEEE, 2013.

18. Barry Boehm. A view of 20th and 21st century software engineering. In *Proceedings of the 28th international conference on Software engineering*, pages 12–29. ACM, 2006.
19. L. C. Briand, J. Wust, S. V. Ikonomovski, and H. Lounis. Investigating quality factors in object-oriented designs: an industrial case study. In *Proceedings of the 1999 International Conference on Software Engineering (IEEE Cat. No.99CB37002)*, pages 345–354, May 1999.
20. Lionel C. Briand, Jürgen Wüst, John W. Daly, and D. Victor Porter. Exploring the relationships between design measures and software quality in object-oriented systems. *Journal of Systems and Software*, 51(3):245 – 273, 2000.
21. Fred P. Brooks, Jr. The mythical man-month. *SIGPLAN Not.*, 10(6):193–, April 1975.
22. Frederick P Brooks. The mythical man-month. *Datamation*, 20(12):44–52, 1974.
23. Frederick P Brooks Jr. *The Mythical Man-Month: Essays on Software Engineering, Anniversary Edition*, 2/E. Pearson Education India, 1995.
24. Andrea Capiluppi, Jesus M Gonzalez Barahona, and Israel Herreraiz. Adapting the “staged model for software evolution” to floss. 2007.
25. Michelle Cartwright and Martin Shepperd. An empirical investigation of an object-oriented software system. *Software Engineering, IEEE Transactions on*, 26:786 – 796, 09 2000.
26. Marcelo Cataldo, Matthew Bass, James D Herbsleb, and Len Bass. On coordination mechanisms in global software development. In *International Conference on Global Software Engineering (ICGSE 2007)*, pages 71–80. IEEE, 2007.
27. Marcelo Cataldo and James D Herbsleb. Coordination breakdowns and their impact on development productivity and software failures. *IEEE Transactions on Software Engineering*, 39(3):343–360, 2013.
28. A. Cockburn and J. Highsmith. Agile software development, the people factor. *Computer*, 34(11):131–133, Nov 2001.
29. Alistair Cockburn. *Agile software development: the cooperative game*. Pearson Education, 2006.
30. Giulio Concas, Michele Marchesi, Alessandro Murgia, and Roberto Tonelli. An empirical study of social networks metrics in object-oriented software. *Adv. Soft. Eng.*, 2010:4:1–4:21, January 2010.
31. Valerio Cosentino, Javier L Cánovas Izquierdo, and Jordi Cabot. A systematic mapping study of software development with github. *IEEE Access*, 5:7173–7192, 2017.
32. Sarah Dahab, Erika Fernanda Silva Balocchi, Stephane Maag, Ana Rosa Cavalli, and Wissam Mallouli. Enhancing software development process quality based on metrics correlation and suggestion. In *ICSOFT 2018: 13th International Conference on Software Technologies*, pages 120–131. Scitepress, 2018.
33. Cleidson RB de Souza, David Redmiles, Li-Te Cheng, David Millen, and John Patterson. Sometimes you need to see through walls: a field study of application programming interfaces. In *Proceedings of the 2004 ACM conference on Computer supported cooperative work*, pages 63–71. ACM, 2004.
34. Luiz Felipe Dias, Igor Steinmacher, Gustavo Pinto, Daniel Alencar da Costa, and Marco Gerosa. How does the shift to github impact project collaboration? In *Software Maintenance and Evolution (ICSME), 2016 IEEE International Conference on*, pages 473–477. IEEE, 2016.
35. K. El Emam, S. Benlarbi, N. Goel, and S. N. Rai. The confounding effect of class size on the validity of object-oriented metrics. *IEEE Transactions on Software Engineering*, 27(7):630–650, July 2001.
36. Karim O. Elish and Mahmoud O. Elish. Predicting defect-prone software modules using support vector machines. *Journal of Systems and Software*, 81(5):649 – 660, 2008. Software Process and Product Measurement.
37. Khaled El Emam, Walcelio Melo, and Javam C. Machado. The prediction of faulty classes using object-oriented design metrics. *Journal of Systems and Software*, 56(1):63 – 75, 2001.
38. Norman Fenton, Martin Neil, William Marsh, Peter Hearty, David Marquez, Paul Krause, and Rajat Mishra. Predicting software defects in varying development lifecycles using bayesian nets. *Information and Software Technology*, 49(1):32 – 43, 2007. Most Cited Journal Articles in Software Engineering - 2000.

39. Brian Fitzgerald and David L Parnas. Making free/open-source software (f/oss) work better. In *Proceedings do Workshop da Conferência XP2003, Genova*. Citeseer, 2003.
40. Aakash Gautam, Saket Vishwasrao, and Francisco Servant. An empirical study of activity, popularity, size, testing, and stability in continuous integration. In *Proceedings of the 14th International Conference on Mining Software Repositories*, pages 495–498. IEEE Press, 2017.
41. T Girba, A Kuhn, M Seeberger, and Stéphane Ducasse. How developers drive software evolution. volume 2005, pages 113– 122, 10 2005.
42. Rebecca E Grinter, James D Herbsleb, and Dewayne E Perry. The geography of coordination: Dealing with distance in r&d work. In *Proceedings of the international ACM SIGGROUP conference on Supporting group work*, pages 306–315. ACM, 1999.
43. Dharmendra Lal Gupta and Kavita Saxena. Software bug prediction using object-oriented metrics. *Sādhanā*, 42(5):655–669, May 2017.
44. Monika Gupta, Ashish Sureka, and Srinivas Padmanabhu. Process mining multiple repositories for software defect resolution from control and organizational perspective. In *Proceedings of the 11th Working Conference on Mining Software Repositories*, pages 122–131. ACM, 2014.
45. Peng He, Bing Li, Xiao Liu, Jun Chen, and Yutao Ma. An empirical study on software defect prediction with a simplified metric set. *Information and Software Technology*, 59:170 – 190, 2015.
46. J. D. Herbsleb and A. Mockus. An empirical study of speed and communication in globally distributed software development. *IEEE Transactions on Software Engineering*, 29(6):481–494, June 2003.
47. James Herbsleb. Socio-technical coordination (keynote). In *Companion Proceedings of the 36th International Conference on Software Engineering*, ICSE Companion 2014, pages 1–1, New York, NY, USA, 2014. ACM.
48. James D Herbsleb and Rebecca E Grinter. Splitting the organization and integrating the code: Conway’s law revisited. In *Proceedings of the 1999 International Conference on Software Engineering (IEEE Cat. No. 99CB37002)*, pages 85–95. IEEE, 1999.
49. Abram Hindle, Daniel M German, and Ric Holt. What do large commits tell us?: a taxonomical study of large commits. In *Proceedings of the 2008 international working conference on Mining software repositories*, pages 99–108. ACM, 2008.
50. Gregory W Hislop, Michael J Lutz, J Fernando Naveda, W Michael McCracken, Nancy R Mead, and Laurie A Williams. Integrating agile practices into software engineering courses. *Computer science education*, 12(3):169–185, 2002.
51. Oskar Jarczyk, Błażej Gruszka, Szymon Jaroszewicz, Leszek Bukowski, and Adam Wierzbicki. Github projects. quality analysis of open-source software. In *International Conference on Social Informatics*, pages 80–94. Springer, 2014.
52. R. Jayanthi and Lilly Florence. Software defect prediction techniques using metrics based on neural network classifier. *Cluster Computing*, Feb 2018.
53. Andrejs Jermakovics, Alberto Sillitti, and Giancarlo Succi. Mining and visualizing developer networks from version control systems. In *Proceedings of the 4th International Workshop on Cooperative and Human Aspects of Software Engineering*, CHASE ’11, pages 24–31, New York, NY, USA, 2011. ACM.
54. Jirayus Jiarpakdee, Chakkrit Tantithamthavorn, and Ahmed E. Hassan. The impact of correlated metrics on defect models, 2018.
55. Marian Jureczko and Lech Madeyski. Towards identifying software project clusters with regard to defect prediction. In *Proceedings of the 6th International Conference on Predictive Models in Software Engineering*, PROMISE ’10, pages 9:1–9:10, New York, NY, USA, 2010. ACM.
56. Eirini Kalliamvakou, Georgios Gousios, Kelly Blincoe, Leif Singer, Daniel M. German, and Daniela Damian. The promises and perils of mining github. In *Proceedings of the 11th Working Conference on Mining Software Repositories*, MSR 2014, pages 92–101, New York, NY, USA, 2014. ACM.
57. Chris F Kemerer and Mark C Paulk. The impact of design and code reviews on software quality: An empirical study based on psp data. *IEEE transactions on software engineering*, 35(4):534–550, 2009.

58. Patrick Knab, Martin Pinzger, and Abraham Bernstein. Predicting defect densities in source code files with decision tree learners. In *Proceedings of the 2006 International Workshop on Mining Software Repositories*, MSR '06, pages 119–125, New York, NY, USA, 2006. ACM.
59. Stefan Koch and Georg Schneider. Effort, co-operation and co-ordination in an open source software project: Gnome. *Information Systems Journal*, 12(1):27–42, 2002.
60. Robert E. Kraut and Lynn A. Streeter. Coordination in software development. *Commun. ACM*, 38(3):69–81, March 1995.
61. Sandeep Krishnamurthy. Cave or community? an empirical examination of 100 mature open source projects (originally published in volume 7, number 6, june 2002). *First Monday*, 0(0), 2005.
62. Chandan Kumar and Dilip Kumar Yadav. Software defects estimation using metrics of early phases of software development life cycle. *International Journal of System Assurance Engineering and Management*, 8(4):2109–2117, Dec 2017.
63. Eetu Kupiainen, Mika V. Mäntylä, and Juha Itkonen. Using metrics in agile and lean software development – a systematic literature review of industrial studies. *Information and Software Technology*, 62:143 – 163, 2015.
64. T. C. Lethbridge. What knowledge is important to a software professional? *Computer*, 33(5):44–50, May 2000.
65. Wei Li and Sallie Henry. Object-oriented metrics that predict maintainability. *Journal of Systems and Software*, 23(2):111 – 122, 1993. Object-Oriented Software.
66. Antonio Lima, Luca Rossi, and Mirco Musolesi. Coding together at scale: Github as a collaborative social network, 2014.
67. J. Ludwig, S. Xu, and F. Webber. Static software metrics for reliability and maintainability. In *2018 IEEE/ACM International Conference on Technical Debt (TechDebt)*, pages 53–54, May 2018.
68. Greg Maday, Vincent Freeh, and Renee Tynan. The open source software development phenomenon: An analysis based on social network theory. *Americas Conference on Information Systems*, 01 1806.
69. Lech Madeyski and Marian Jureczko. Which process metrics can significantly improve defect prediction models? an empirical study. *Software Quality Journal*, 23(3):393–422, Sep 2015.
70. Anirban Majumder, Samik Datta, and K.V.M. Naidu. Capacitated team formation problem on social networks. *Proceedings of the 18th ACM SIGKDD international conference on Knowledge discovery and data mining - KDD '12*, 2012.
71. Shane McIntosh, Yasutaka Kamei, Bram Adams, and Ahmed E. Hassan. An empirical study of the impact of modern code review practices on software quality. *Empirical Software Engineering*, 21(5):2146–2189, Oct 2016.
72. Andrew Meneely, Laurie A. Williams, Will Snipes, and Jason A. Osborne. Predicting failures with developer networks and social network analysis. In *SIGSOFT FSE*, 2008.
73. Tim Menzies, Zach Milton, Burak Turhan, Bojan Cukic, Yue Jiang, and Ayşe Bener. Defect prediction from static code features: current results, limitations, new approaches. *Automated Software Engineering*, 17(4):375–407, Dec 2010.
74. Nikolaos Mittas and Lefteris Angelis. Ranking and clustering software cost estimation models through a multiple comparisons algorithm. *IEEE Trans. Software Eng.*, 39(4):537–551, 2013.
75. Audris Mockus, Roy T Fielding, and James D Herbsleb. Two case studies of open source software development: Apache and mozilla. *ACM Transactions on Software Engineering and Methodology (TOSEM)*, 11(3):309–346, 2002.
76. ABM Moniruzzaman and Dr Syed Akhter Hossain. Comparative study on agile software development methodologies. *arXiv preprint arXiv:1307.3356*, 2013.
77. Stefan Morcov. Complex it projects in education: The challenge. *International Journal of Computer Science Research and Application*, 2:115–125, 2012.
78. Nuthan Munaiah, Steven Kroh, Craig Cabrey, and Meiyappan Nagappan. Curating github for engineered software projects. *Empirical Software Engineering*, 22(6):3219–3253, Dec 2017.
79. Nachiappan Nagappan, Thomas Ball, and Andreas Zeller. Mining metrics to predict component failures. In *Proceedings of the 28th International Conference on Software Engineering*, ICSE '06, pages 452–461, New York, NY, USA, 2006. ACM.

80. Ahmet Okutan and Olcay Taner Yıldız. Software defect prediction using bayesian networks. *Empirical Software Engineering*, 19(1):154–181, Feb 2014.
81. T. J. Ostrand, E. J. Weyuker, and R. M. Bell. Predicting the location and number of faults in large software systems. *IEEE Transactions on Software Engineering*, 31(4):340–355, April 2005.
82. Thomas J. Ostrand, Elaine J. Weyuker, and Robert M. Bell. Programmer-based fault prediction. In *Proceedings of the 6th International Conference on Predictive Models in Software Engineering*, PROMISE ’10, pages 19:1–19:10, New York, NY, USA, 2010. ACM.
83. Kevin Peterson. The github open source development process. 12 2013.
84. Martin Pinzger, Nachiappan Nagappan, and Brendan Murphy. Can developer-module networks predict failures? In *SIGSOFT FSE*, 2008.
85. Karuna Prasad, MG Divya, and N Mangala. Statistical analysis of metrics for software quality improvement, 2018.
86. Manjula.C.M. Prasad, Lilly Florence, and Arti Arya. A study on software metrics based software defect prediction using data mining and machine learning techniques. *International Journal of Database Theory and Application*, 7:179–190, 06 2015.
87. Ayushi Rastogi, Nachiappan Nagappan, and Pankaj Jalote. Empirical analyses of software contributor productivity. 2017.
88. Jacek Ratzinger, Thomas Sigmund, and Harald C. Gall. On the relation of refactorings and software defect prediction. In *Proceedings of the 2008 International Working Conference on Mining Software Repositories*, MSR ’08, pages 35–38, New York, NY, USA, 2008. ACM.
89. T. Ravi Kumar, T. Srinivasa Rao, and Sandhya Bathini. A predictive approach to estimate software defects density using weighted artificial neural networks for the given software metrics. In Suresh Chandra Satapathy, Vikrant Bhateja, and Swagatam Das, editors, *Smart Intelligent Computing and Applications*, pages 449–457, Singapore, 2019. Springer Singapore.
90. Arturo Reyes López. Analyzing github as a collaborative software development platform: A systematic review. 2017.
91. Filippo Ricca and Alessandro Marchetto. Are heroes common in floss projects? In *Proceedings of the 2010 ACM-IEEE International Symposium on Empirical Software Engineering and Measurement*, page 55. ACM, 2010.
92. Gregorio Robles and Jesus M Gonzalez-Barahona. Contributor turnover in libre software projects. In *IFIP International Conference on Open Source Systems*, pages 273–286. Springer, 2006.
93. Gregorio Robles, Jesus M Gonzalez-Barahona, and Israel Herraiz. Evolution of the core team of developers in libre software projects. In *Mining Software Repositories, 2009. MSR’09. 6th IEEE International Working Conference on*, pages 167–170. IEEE, 2009.
94. Christoffer Rosen, Ben Grawi, and Emad Shihab. Commit guru: analytics and risk prediction of software commits. In *Proceedings of the 2015 10th Joint Meeting on Foundations of Software Engineering*, pages 966–969. ACM, 2015.
95. Y. Shin, A. Meneely, L. Williams, and J. A. Osborne. Evaluating complexity, code churn, and developer activity metrics as indicators of software vulnerabilities. *IEEE Transactions on Software Engineering*, 37(6):772–787, Nov 2011.
96. Q. Song, Z. Jia, M. Shepperd, S. Ying, and J. Liu. A general software defect-proneness prediction framework. *IEEE Transactions on Software Engineering*, 37(3):356–370, May 2011.
97. R. Subramanyam and M. S. Krishnan. Empirical analysis of ck metrics for object-oriented design complexity: implications for software defects. *IEEE Transactions on Software Engineering*, 29(4):297–310, April 2003.
98. Jessica Suzuki and Edna Dias Canedo. Interaction design process oriented by metrics. In Constantine Stephanidis, editor, *HCI International 2018 – Posters’ Extended Abstracts*, pages 290–297, Cham, 2018. Springer International Publishing.
99. Debbie Tesch, Marion G. Sobol, Gary Klein, and James J. Jiang. User and developer common knowledge: Effect on the success of information system development projects. *International Journal of Project Management*, 27(7):657 – 664, 2009.

100. Venkata Udaya B. Challagulla, Farokh B. Bastani, I-Ling Yen, and Raymond A. Paul. Empirical assessment of machine learning based software defect prediction techniques. *International Journal on Artificial Intelligence Tools*, 17:389–400, 04 2008.
101. Bogdan Vasilescu, Yue Yu, Huaimin Wang, Premkumar Devanbu, and Vladimir Filkov. Quality and productivity outcomes relating to continuous integration in github. In *Proceedings of the 2015 10th Joint Meeting on Foundations of Software Engineering*, pages 805–816. ACM, 2015.
102. T John Vijay, Dr M Gopi Chand, and Dr Harika Done. Software quality metrics in quality assurance to study the impact of external factors related to time. *International Journal of Advanced Research in Computer Science and Software Engineering*, 7(1), 2017.
103. Elaine J. Weyuker, Thomas J. Ostrand, and Robert M. Bell. Do too many cooks spoil the broth? using the number of developers to enhance defect prediction models. *Empirical Software Engineering*, 13(5):539–559, Oct 2008.
104. J. Whitehead. Collaboration in software engineering: A roadmap. In *Future of Software Engineering (FOSE '07)*, pages 214–225, May 2007.
105. Chadd Williams and Jaime Spacco. Szz revisited: verifying when changes induce fixes. In *Proceedings of the 2008 workshop on Defects in large software systems*, pages 32–36. ACM, 2008.
106. Timo Wolf, Adrian Schroter, Daniela Damian, and Thanh Nguyen. Predicting build failures using social network analysis on developer communication. In *Proceedings of the 31st International Conference on Software Engineering*, pages 1–11. IEEE Computer Society, 2009.
107. Timo Wolf, Adrian Schroter, Daniela Damian, and Thanh Nguyen. Predicting build failures using social network analysis on developer communication. In *Proceedings of the 31st International Conference on Software Engineering*, ICSE '09, pages 1–11, Washington, DC, USA, 2009. IEEE Computer Society.
108. Trevor Wood-Harper and Bob Wood. Multiview as social informatics in action: past, present and future. *Information Technology & People*, 18(1):26–32, 2005.
109. Michalis Xenos. Software metrics and measurements. *Encyclopedia of E-Commerce, E-Government and Mobile Commerce*, pages 1029–1036, 01 2006.
110. F. Zhang, A. E. Hassan, S. McIntosh, and Y. Zou. The use of summation to aggregate software metrics hinders the performance of defect prediction models. *IEEE Transactions on Software Engineering*, 43(5):476–491, May 2017.
111. Feng Zhang, Audris Mockus, Iman Keivanloo, and Ying Zou. Towards building a universal defect prediction model. In *Proceedings of the 11th Working Conference on Mining Software Repositories*, MSR 2014, pages 182–191, New York, NY, USA, 2014. ACM.
112. Thomas Zimmermann, Nachiappan Nagappan, Harald Gall, Emanuel Giger, and Brendan Murphy. Cross-project defect prediction: A large scale experiment on data vs. domain vs. process. In *Proceedings of the the 7th Joint Meeting of the European Software Engineering Conference and the ACM SIGSOFT Symposium on The Foundations of Software Engineering*, ESEC/FSE '09, pages 91–100, New York, NY, USA, 2009. ACM.
113. Thomas Zimmermann, Rahul Premraj, and Andreas Zeller. Predicting defects for eclipse. In *Proceedings of the Third International Workshop on Predictor Models in Software Engineering*, PROMISE '07, pages 9–, Washington, DC, USA, 2007. IEEE Computer Society.